

/hai — Week 4 Report

Sai Pranavi Jeedigunta | Mycroft Project 29

July 20 – July 26, 2026 (*report compiled July 24, mid-week*)

Hours Log

Date	Activity	Hours
Mon July 20	Sprint call with Shradha (45 min); drafted Substack article	~3.5
Tue July 21	Completed and published Substack article; obtained GitHub access to Project 29 repo	~3.0
Wed-Thu July 22-23	Completed Mycroft book Chapters 1 (The Fluency Trap) and 2 (The Reallocation Principle) in full — reading plus all conceptual and “Your Own Mycroft” exercises, including a full claim-audit and effort-plan for a live investment thesis (Titan Company)	~5.5
Fri July 24	Mandatory Prof. Brown office hours (video report transition); Group A/B pipeline hardening — local DB setup, root-cause diagnosis, parameterized query fix, empty-content filter fix, dedup logic, HTML escaping, config fixes, live verification on 370-item run	~8.0
Total		~20.0

Accomplishments

1. Published second Mycroft Substack article: “The Classification Problem Starts Before the Keywords Run”

Original analysis, not a summary — argues the pipeline's classification failures are architectural (ingestion-layer) rather than fixable through keyword tuning, and pre-empts the “add an LLM” solution by explaining why an LLM inherits upstream data corruption rather than fixing it. Built directly on primary research from the Week 3 node-by-node audit.

2. Completed first two chapters of Mycroft book curriculum, including all exercises

Completed Chapter 1 (The Fluency Trap) and Chapter 2 (The Reallocation Principle) in full — reading plus all conceptual exercises and all “Your Own Mycroft” project exercises for both chapters. Confirmed with PM that chapters replace the Gru/SDD requirement for this project — closing an item unresolved since Week 1.

Built a full claim-audit (theses/TITAN.md) for a live investment thesis on Titan Company, an India-listed jewelry/watch retailer. Identified that the thesis's load-bearing connective claim — that rising durable-goods spending extends specifically to branded jewelry/watches — is not supported by the underlying survey data (which names vehicles, appliances, and phones, but not jewelry or watches) and correctly withheld “verified” status from that claim rather than letting adjacent verified claims lend it false credibility. Followed with a research-effort plan (effort-plan.md) that reprioritized based on personal decision-impact weighting rather than model-perceived salience — notably, Claim 2 (middle-class growth sourcing) was promoted in priority once real risk weighting was supplied, reversing the AI's initial neutral ranking.

3. Executed and live-verified a full round of production pipeline hardening

Stood up local PostgreSQL environment, created a tracked working copy (workflow.dev.json), and traced the root-cause chain from a remote DB timeout to a VARCHAR overflow. Completed and verified on a real 370-item

run: parameterized SQL insert (eliminating injection risk), recovery of 73 previously silently-dropped SEC/exchange filings, robust duplicate detection, HTML escaping, corrected file path and timezone, feed-failure isolation config, and threshold alignment across report and email nodes. Deliberately held Node 10's urgency-scoring bug unfixed to preserve it as the Layer-2 measurement baseline — documented in FINDINGS.md.

4. Attended mandatory office hours announcing a new program-wide renewal requirement, and completed it the same reporting period

At Prof. Brown's mandatory Friday office hours (July 24), a video report was announced as a newly-implemented requirement for OPT renewal, applying program-wide rather than to this project specifically. A video report ("The Pipeline That Was Lying to Me") was built and completed the following Monday, documenting the empty-content filter bug (Node 9) with narration fact-checked line-by-line against live pipeline output (FINDINGS.md). Full build documentation (README, BUILD-LOG, FACTCHECK, SOURCES, PEDAGOGY, SHOTLIST) was produced and pushed to github.com/nikbearbrown/humanitarians-youtube, branch `sai-pranavi-j/2026-07-26-pipeline-lying-report`. This satisfies the new requirement ahead of renewal.

Struggle Documentation

The central difficulty this week was sequencing, not execution: distinguishing fixes that could go in immediately from ones that would corrupt future measurement if fixed too early. The Group C decision — leaving Node 10's baseline-5 compression untouched — required holding two competing instincts at once: the instinct to fix every visible bug, and the recognition that fixing this one now would erase the noise-baseline needed to prove later work (Layer 2/3) actually improved anything. This was resolved by writing the misfire evidence (Medicare→10/Critical, Nasdaq cluster→9/Critical) into FINDINGS.md before touching anything else, so the baseline exists on record independent of the fix decision.

A secondary struggle: reading two chapters on AI-assisted judgment (the fluency trap, the reallocation principle) while mid-week in actual production debugging created a live test of the material — several of Chapter 1's supervision questions (scope, approval, verification) map almost exactly onto decisions made in the Group A/B triage, which was not planned but became apparent while writing this report.

Tier Justification

Tier 5 — Causal Reasoning: The Node 10 hold is a causal argument, not a bug-fix decision: fixing the urgency-scoring compression now would destroy the ability to measure whether later architectural work (Layer 2/3) actually caused an improvement. Distinguishing "bug to fix" from "baseline to preserve" required reasoning about what the fix would do to a future measurement, not just what it does to current output.

Tier 4 — Judgment and Supervision: Root-cause diagnosis (remote timeout → local repoint → VARCHAR overflow) and the Group A/B/C triage itself required judgment about sequencing — recognizing that some fixes are safe regardless of order (reliability) while others change what the system detects and must be sequenced deliberately (detection-behavior).

Tier 2 — Embodied: Live verification against a real 370-item production run — reading actual output distributions, confirming 46 new / 324 dedup / 0 errors — is knowledge built by running the system, not by reading its code.

Tier 4/5 — Judgment Transfer: The Titan Company claim-audit demonstrates the same judgment skill from the Project 29 pipeline work transferring to an entirely unrelated domain (personal investment research) with no overlap in subject matter. Catching that a survey's named categories (durable goods) do not actually include the thesis's specific claim (jewelry/watches) is the same discipline as catching that Node 8 mislabels non-financial Federal Register items as Securities: recognizing that adjacent, verified facts do not automatically lend credibility

to an unverified connective claim sitting between them.

Tier 1 — AI Fluency: Completed Chapters 1–2 of the Mycroft curriculum in full, which directly names the failure mode this week's Substack article and pipeline work both independently surfaced (fluent output outrunning its evidence). The Chapter 2 phase-gate exercise additionally surfaced a genuine design gap (undefined ownership for adjusting a threshold on a customer with a known payment plan) through self-review of a completed design — itself an instance of the verification discipline the chapter teaches.

Unresolved

- Group A: item A7 (scope email-sent UPDATE) not started
- Group B: B2 (source classification), B3 (URL unwrap) not started
- Group D: structured run logging, failure monitoring — not started

Next Steps — Week 5

- Attend and act on outcomes from July 24 office hours (video report format/cadence)
- Complete Chapter 1 and 2 exercises before starting Chapter 3
- Resume Group B/D fixes: B2 source classification first (flagged as highest-impact remaining gap)
- Verify unemployment-days figure directly rather than carrying forward
- Update Google Sheet with Week 4 Substack link and GitHub branch (done — see prior summary)

Compliance Notes

Item	Status
Volunteer agreement	Signed and countersigned — June 29 to July 31, 2026
SEVP portal	Updated — Humanitarians AI Inc. registered June 29, 2026
Unique ID	241
Unemployment days used	Confirmed by volunteer: approximately 11 of 90 allowed
Gru/SDD	Not required for this project — chapters substitute (confirmed with PM, Week 4)
CRITIQ	Confirmed standing requirement — all published Substack articles cleared through CRITIQ prior to publication