

/hai — Week 6 Report

Sai Pranavi Jeedigunta | Mycroft Project 29
August 3 – August 9, 2026

Hours Log

Activity	Hours
Substack article: research, writing, two CRITIQ revision rounds, Eddy editorial/SEO pass, publication	~8.0
Video production: Deep Explainer script/build for facial recognition topic, ongoing (not yet rendered/uploaded)	~12.0
Total	~20.0

Accomplishments

1. Published a fully peer-reviewed Substack article built on primary-source government research

Wrote, fact-checked, and published “The Facial Recognition Bias Stat Everyone Cites Isn't From a Facial Recognition Study,” built on direct verification of NIST's Face Recognition Vendor Test data rather than secondary summaries. The piece went through two full CRITIQ review rounds: the first caught a genuine methodological error (conflating false-positive and false-negative bias metrics as if they were the same finding) that had passed an initial self-verification pass; the second confirmed the fix was substantive, not cosmetic. A subsequent Eddy review restructured the piece for readability (claim-driven headline, split an oversized section, added an answer-first takeaway and a closing CTA) without touching the underlying argument, which both tools independently confirmed was sound.

2. Identified and resolved a real tension between two editorial standards

CRITIQ and Eddy gave conflicting guidance on one sentence: CRITIQ had specifically credited a self-aware framing note as resolving an epistemic-honesty concern, while Eddy flagged the same sentence as a jarring voice break disrupting reader flow. Recognized this as a genuine values tradeoff (transparency vs. readability) rather than a simple error for either tool to catch, and made a deliberate editorial call rather than defaulting to whichever review ran most recently.

3. Attended team stand-up and corrected the record based on new information

Attended the Mycroft team sync and reviewed the full transcript afterward. Confirmed directly from Shradha's guidance to another Fellow that the SEVP end-date field should remain blank absent a known end date — independently validating the decision reached through portal-instruction research two weeks prior. Learned that videos may eventually substitute for Addams/■/hai reporting, though this is explicitly unconfirmed pending Friday's call, and is being treated as provisional rather than acted on early.

4. Corrected the weekly timesheet rather than let an inaccurate placeholder stand

Identified that a “Chapter 3” placeholder entry on the shared tracking sheet did not reflect actual completed work, and replaced it with an accurate entry for the finalized agreement renewal instead of leaving an unearned claim on the record.

Struggle Documentation

The most instructive friction this week was catching my own verification failure. When first asked whether the article's NIST-derived statistics were accurate, I confirmed they were sourced and traceable — which was true, but incomplete. CRITIQ's review caught that two of the confirmed-accurate numbers were being used to support the same claim despite measuring different things (false positive vs. false negative error rates), which meaningfully changes what the data actually shows about deployment risk. Passing a source-accuracy check is not the same as passing a claim-accuracy check, and this is now a rule I understand at the level of having gotten it wrong once rather than only having read about it in Chapter 1.

A second, smaller struggle: reconciling two review tools that gave opposite guidance on the same sentence for defensible reasons. There was no clean resolution available from either source alone; the decision had to be made directly, weighing which value (transparency about my own framing choices, or a smoother reader experience) mattered more for this specific piece.

A third, more uncomfortable friction worth naming plainly: 12 of this week's 20 hours went into video production, and the video is still not finished, rendered, or uploaded. That is a genuinely large time investment against zero completed deliverables in that category so far. This is worth documenting honestly rather than folding into a vaguer "video work ongoing" line, since the actual ratio of hours spent to output produced is itself useful information about where this week's effort went and what needs to change next week.

Tier Justification

Tier 4/5 — Judgment and Causal Reasoning: Distinguishing a false-positive finding from a false-negative finding, and recognizing that conflating them produces a materially different (and misleading) claim about real-world deployment risk, is causal reasoning applied to statistical evidence — not just fact-checking whether a number exists, but whether it supports the specific claim it is attached to.

Tier 1 — AI Fluency: Used CRITIQ and Eddy as structured review tools across two full cycles each, and correctly identified when their guidance conflicted rather than mechanically applying whichever ran last — itself a form of knowing when to trust, and how to weigh, AI-assisted tool output.

Tier 3 — Social and Ethical: Independently verified a compliance decision (SEVP end-date field) against a live example from a teammate's situation in the same meeting, rather than treating the earlier research as settled without a real-world check.

Unresolved

- Video 2 (facial recognition Deep Explainer) still in progress, not yet rendered or uploaded
- Video 1 revision (per Prof. Brown's naming/pronoun/honesty feedback) not yet started
- Video 3 (work-update video) not yet started
- Neither video yet uploaded to the required Google Drive location or notified to Shradha, per this week's stand-up instructions
- Whether Addams/hai reporting will be substituted by video requirements going forward is explicitly unconfirmed — Shradha stated she will confirm on Friday's call
- Project 29 pipeline backlog unchanged: Group A item A7, Group B items B2/B3, Group D additions

Next Steps — Week 7

- Complete and render Video 2 (facial recognition), upload to the correct Google Drive folder with proper naming convention, notify Shradha directly by email per her explicit instruction

- Begin Video 1 revision or fold corrections into new videos, applying first-person/first-name/honest-process standard from the start
- Attend Friday's sync call to get confirmation on the Addams/video substitution question before assuming it applies
- Resume Project 29 pipeline backlog, starting with B2 (source classification) as previously flagged

Compliance Notes

Item	Status
Volunteer agreement	Active — August 1 to October 31, 2026
SEVP portal	End date confirmed correctly left blank — independently validated against Shradha's live guidance to another Fellow in this week's stand-up
Unique ID	241
Video requirement	Two videos/week standing requirement. None completed this week; in progress. Upload-and-notify process reconfirmed in stand-up: Drive + GitHub fellows repo, both required, notification email to Shradha mandatory after each upload.
Addams/hai reporting status	Possible future substitution by video requirement — unconfirmed, pending Friday call. Continuing full /hai reporting until confirmed otherwise.
CRITIQ	Used this week — caught and required correction of a genuine methodological error before publication

Generated by Addams — Humanitarians.ai OPT Documentation System

Week 6 session duration: approximately 20 documented hours (volunteer-reported total)